

veeAMON2023

Veeam Repository Sizing & Best Practices



Adam Congdon
Solutions Architect



Olivier Rossi
Senior Solutions Architect




Agenda

- Backup repository
- Scale-out Backup Repository (SOBR)
- Object storage repository
- NAS backup & cache repository



Main components



Veeam Backup & Replication™

1. Veeam® backup server
2. Database server(s)
3. Proxy server
4. Repository server
5. Veeam Backup Enterprise Manager
6. Combo server (proxy and repository)



Backup repository





Repository server

How to size: compute



Per task (VM or agent)

One CPU core for every task
(usually one-third of proxy cores)

Plus, one CPU core for every concurrent offload task



Per task (VM or agent)

4GB for every task (min 2GB)

+2GB for every concurrent offload task





Repository server




How to size: capacity

Space usage



Calculated based off retention points and backup method used

See calculator.veeam.com/vbr for calculations



Repository server

Input values:

300TB source data

Retention

- On-site: seven days (synth full Saturdays), ReFS.
- Off-site: four weeks/six months/four years.

Offload to object after seven days

Daily change rate 10%



Example calculation

On-site:

$$(\text{full}) + (\text{inc.} * (\text{days} - 1 + 6)) + \text{workspace} = \text{capacity}$$

$$\frac{300\text{TB}}{50\% \text{ data reduction}} = \mathbf{150\text{TB (full)}}$$

$$150\text{TB (full)} * 10\% = \mathbf{15\text{TB (incremental)}}$$

$$(150) + (15 * (7 - 1 + 6)) + \mathbf{53.1} = \mathbf{383.1\text{TB}}$$

Off-site (object storage):

$$(150) + (15 * (4\text{w} + 6\text{m})) + (4 * 150) = \mathbf{930\text{TB Min}^*}$$

veeAMON2023



Scale-out Backup Repository™





Considerations

1. Try to keep extent count < 16
2. Pay attention to placement policy
3. Separate VM/Agent backups from NAS/plugin backups
4. Metadata extents for NAS
5. Cloud offload
6. Block cloning does not span extents





NAS & metadata extents

1. When mixing dedupe storage with SSD storage: set SSD storage as metadata extents to improve transform performance
2. Assign via PowerShell
3. Metadata replica will exist on data extent also — not used

veeAMON2023



Object storage repository





Direct to object



Considerations

1. Gateway(s)
2. Bucket security
3. Bandwidth
4. Object store limitations
5. Storage class (minimum stay)
6. Egress cost (Hotel California Syndrome)
7. Capacity, API cost, immutability





Direct to object

Gateway(s) & bucket security



Connection Type

Choose a connection mode for this object storage repository. For the Direct mode, ensure backup proxies and backup agents have direct network access to object storage. For the Gateway mode, we recommend having at least two gateway servers for redundancy.

☐ Direct
Backup proxies and agents will connect directly to object storage.

☒ Through a gateway server
Backup server will automatically determine the most suitable gateway server from the following list:

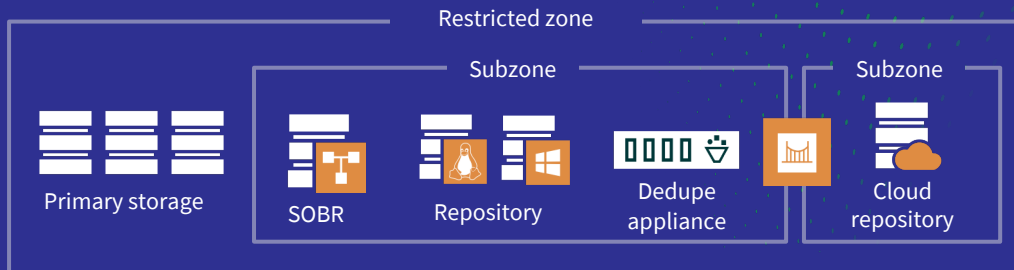
Name
<input type="checkbox"/> cacherp01.protekt.local
<input type="checkbox"/> ec2-52-41-47-75-us-west-2.compute.amazonaws.com
<input checked="" type="checkbox"/> fp01.protekt.local
<input type="checkbox"/> fp02.protekt.local
<input type="checkbox"/> gnu01.protekt.local
<input type="checkbox"/> helper01.protekt.local
<input type="checkbox"/> hki01.protekt.local
<input type="checkbox"/> hki02.protekt.local
<input type="checkbox"/> img01.protekt.local
<input type="checkbox"/> vbr12.protekt.local
<input type="checkbox"/> vsm01.protekt.local

Connection mode:
fp01.protekt.local

Specify whether object storage should be accessed directly or via selected gateway servers.

OK Cancel

```
EDIT POLICY
8      "s3:PutObject",
9      "s3:GetObject",
10     "s3:DeleteObject",
11     "s3:GetBucketLocation",
12     "s3:GetBucketVersioning",
13     "s3:GetBucketObjectLockConfiguration"
14   ],
15   "Resource": [
16     "arn:aws:s3:::orossisecurebucket01/*",
17     "arn:aws:s3:::orossisecurebucket01"
18   ]
```





Direct to object

Bandwidth


B/W = Backup data/backup window







Object store limitations

Object metadata implementation
(DB, on-disk...)



New Scale-out Backup Repository

 **Performance Tier**
Select backup repositories to use as the landing zone and for the short-term retention.

Name	Extents:						
Performance Tier	<table border="1"><thead><tr><th>Name</th></tr></thead><tbody><tr><td> mbucket01</td></tr><tr><td> mbucket02</td></tr><tr><td> </td></tr><tr><td> </td></tr><tr><td> </td></tr></tbody></table>	Name	 mbucket01	 mbucket02			
Name							
 mbucket01							
 mbucket02							
Capacity Tier							
Summary							



Direct to object

Storage class

	S3 Standard	S3 Intelligent-Tiering*	S3 Standard-IA	S3 One Zone-IA†	S3 Glacier	S3 Glacier Flexible Retrieval	S3 Glacier Deep Archive
					Instant Retrieval		
Designed for durability	100.00%	100.00%	100.00%	100.00%	100.00%	100.00%	100.00%
	(11 9's)	(11 9's)	(11 9's)	(11 9's)	(11 9's)	(11 9's)	(11 9's)
Designed for availability	99.99%	99.90%	99.90%	99.50%	99.90%	99.99%	99.99%
Availability SLA	99.90%	99%	99%	99%	99%	99.00%	99.90%
Availability Zones	≥3	≥3	≥3	1	≥3	≥3	≥3
Minimum capacity charge per object	N/A	N/A	128 KB	128 KB	128 KB	40 KB	40 KB

Minimum storage duration charge

Retrieval charge

First byte latency

Storage type

Lifecycle transitions

Cool and Archive

In addition to the per object charge, there is a retrieval charge of 180 days. Additionally, there is a deletion period of 30 days. This charge is prorated. For example, if a blob is moved to the Archive tier and then deleted or moved to the Hot tier after 45 days, the customer is charged an early deletion fee for 135 (180 minus 45) days of storage in the Archive tier.

☒ Store archived backups as standalone fulls

To reduce storage costs, we store each backup as a delta from the previous one. For the long-term archive reliability considerations, you may instead choose to store each backup as a standalone full, without any dependencies on previous backups.

☒ Archive backups only if the remaining retention time is above minimal storage period

Avoid processing costs and penalties associated with early data deletion by not archiving backups which will have to be deleted before the minimum storage period expires.
Minimum storage duration for the selected repository: 90 days

OK

Cancel

180 days

per GB retrieved

hours

Object

Yes

period of early



Direct to object

Egress cost

Data Transfer OUT From Amazon S3 To Internet

AWS customers receive 100GB of data transfer out to the internet free each month, aggregated across all AWS Services and Regions (except China and GovCloud). The 100 GB free tier for data transfer out to the internet is global and does not apply separately or individually to AWS Regions.

First 10 TB / Month	\$0.09 per GB
Next 40 TB / Month	\$0.085 per GB
Next 100 TB / Month	\$0.07 per GB
Greater than 150 TB / Month	\$0.05 per GB

Internet Egress (routed via Microsoft Premium Global Network)

Source Continent	First 100GB / Month	Next 10TB / Month	Next 40TB / Month	Next 100TB / Month	Next 350TB / Month	
From North America, Europe to any destination	Free	\$0.087 per GB	\$0.083 per GB	\$0.07 per GB	\$0.05 per GB	\$0.05 per GB
From Asia (China excluded), Australia, MEA to any destination	Free	\$0.12 per GB	\$0.085 per GB	\$0.082 per GB	\$0.08 per GB	\$0.08 per GB
From South America to any destination	Free	\$0.181 per GB	\$0.175 per GB	\$0.17 per GB	\$0.16 per GB	\$0.16 per GB




Direct to object

How to size: compute

- Same math as backup repository
- Apply to gateway(s) server(s)

How to size: capacity





Direct to object


Input values:

300TB source data

Retention: seven days, four weeks

Daily change rate 10%

Example calculation



SOBR with Capacity & Archive Tier

Example scenario:


Retention: seven days, four weeks, six months, four years

Capacity offload: after seven days

Archive offload: after 180 days (standalone fulls)

Performance Tier: leverages block clone

Backup method: forward incremental (i.e., synthetic fulls)



API calls

Direct to object

of PUTs = source data / 1MB

Express per month


First month: full PUTs + 30 inc. PUTs

Last month: 31 inc. PUTs

Immutability

Immutability PUTs(10 days) = FULL/1MB

SOBR capacity and archive tiers





NAS backup & cache repository



NAS repository



Cache repository



CPU

Two-thirds of file proxy cores



RAM

4GB per CPU core



Disk

~1GB or 5% of source data (if direct to object)

Repository/gateway server



CPU

1/2 of file proxy cores



RAM

4GB per CPU core



NAS repository

Repository capacity

Typical data reduction: [100-70%]

Typical change rate: [1-3%]

Metadata + workspace for disk repo: 20% * (full + inc.)

Metadata for object repo: 5% * source data

Repository capacity = full + incremental + meta

API calls

of PUTs_(data) = Full(or inc.)/64MB

of PUTs_(metadata) = 20% of PUTs_(data)





Thank you!

veeAMON2023

